

# Computer Vision for HCI

Image Segmentation  
and  
Template Matching

# Image Segmentation

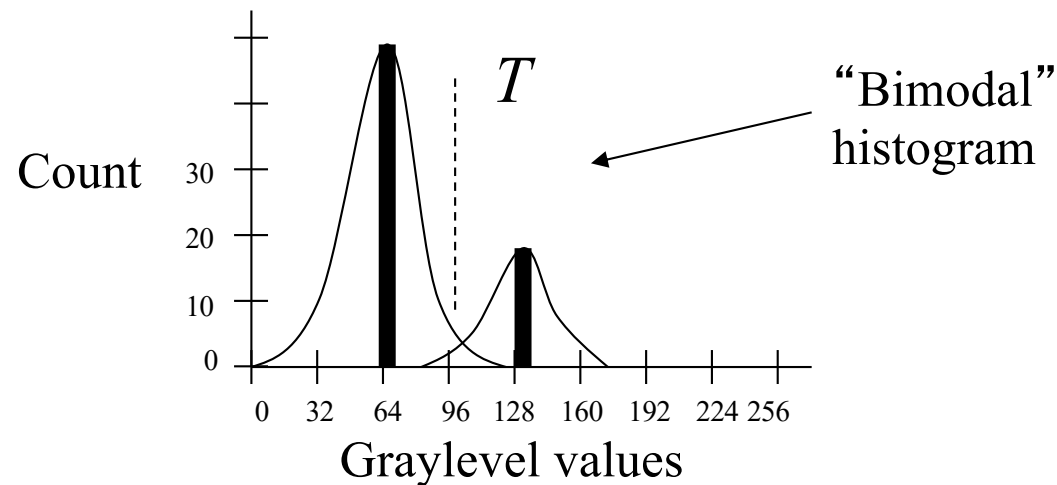
- Goal: Partition an image into distinct regions containing each pixels with similar attributes
- Use “*discontinuity*” or “*similarity*” approach
  - *Discontinuity*: segment into regions based on discontinuity (gradient or edge detection)
  - *Similarity*: Merge similar regions (clustering, region growing etc.)
- Topics
  - Simple Segmentation
  - Segmentation by Clustering
  - Superpixel Segmentation

# Goal



# Recap: Otsu's Simple Segmentation

- Distribution of graylevels can be used to determine binary threshold
- Histogram graphs number of pixels in the image with a particular graylevel, as a function of the possible graylevels
  - Find peaks and set threshold between peaks



# Otsu's Method


- “A threshold selection method from graylevel histograms”, IEEE Trans on Sys., Man, and Cyb., Vol 9, No 1, pp 62-66, 1979.
  - Basic idea: threshold is chosen such that the division in the histogram yields the largest reduction in standard deviation of the pixel intensities (black, white)
  - Matlab: `graythresh()`



# Image Segmentation by Clustering

Identify groups of pixels that “go together”

Each “point” is a pixel in  
color space (3-D: RGB)



- K-Means
- Mean-Shift Clustering

# K-Means

# K-means Clustering

- Each “point” is a 3-D vector of color (RGB)
- Initialization:
  - Choose  $k$  cluster centers (how pick  $k$ ?)
- Repeat:
  - **Assignment step:**
    - For every point, find its closest center
  - **Update step:**
    - Update every center as the mean of its assigned points
- Until:
  - The maximum number of iterations is reached, or
  - No changes during the assignment step, or
  - The average distortion per point drops very little



# K-means: Initialization

- K-means is *extremely sensitive* to initialization
- Bad initialization can lead to
  - Poor convergence speed
  - Poor overall clustering
- How to initialize?
  - Randomly from data
  - Try to find K “spread-out” points
- Try multiple initializations and pick best result
  - Minimize total “distortion” (sum of distances of points from their cluster centers)

$$J(\mu, r) = \sum_{n=1}^N \sum_{k=1}^K \delta_{nk} ||x_n - \mu_k||^2$$

Cluster center    Data

Whether  $x_j$  is assigned to  $\mu_i$

# Example

Original Image



Segmented Image when  $K = 3$






# Mean-Shift

# Mean-Shift segmentation

- Recall Mean-Shift tracking lecture...
- Used here for unsupervised clustering
  - Unlike K-means, do not need initial ‘K’
- Assigns the data points to clusters iteratively by shifting points towards the local modes
  - Mode: The highest density of data points in the region, in the context of the Mean-Shift



# Mean-Shift clustering

- **Relate to “Kernel Density Estimation”**
  - Imagine the data sampled from a probability distribution
  - Estimate the underlying distribution (also called the probability density function) for a set of data
    - Place kernel on each point (think weighing function)
    - Add all the individual kernels generates a probability surface (e.g., density function)



# Mean-Shift Clustering

- Idea:
  - Make points climb up the hill to the nearest peak on the density surface
  - *Iteratively* shift each point uphill until it reaches a peak



# Mean-Shift Algorithm

- Define  $x$ 
  - Color only:  $[R, G, B]$
  - Spatial and Color:  $[x\text{-loc}, y\text{-loc}, R, G, B]$
- For each datapoint  $x$ , find the neighboring points  $N(x)$  of  $x$ , given Kernel function/window  $K$
- For each  $x$ , calculate the ***mean shift***  $m(x)$ :


$$m(x) = \frac{\sum_{x_i \in N(x)} K(x_i - x) x_i}{\sum_{x_i \in N(x)} K(x_i - x)}$$

- Then update each  $x$  with  $x \leftarrow m(x)$
- Repeat  $n$  times or until the points stabilize


More details:

<https://spin.atomicobject.com/2015/05/26/mean-shift-clustering/>






(a)




(b)





# RGB Visualization



# Examples



# Superpixels



# Supervoxel Segmentation

- “Supervoxels” capture local visual redundancy in the image
  - **SLIC** Supervoxel algorithm



# SLIC Superpixel Segmentation

- Generated by clustering pixels based on:
  - Color similarity, and
  - Spatial proximity in the image
- Employ 5-D vector per pixel: [lab, x, y]
  - lab = pixel color vector in CIELAB color space
    - (l): intensity, (a, b): color
    - Perceptual color space with Euclidean distance properties
  - x, y = pixel position (or row, col)

# SLIC Supersixel Algorithm

- Initially choose  $K$  = number of desired supersixels
- Divide image into regular “grid” steps  $S$  (for the  $K$ )

---

**Algorithm 1** Efficient supersixel segmentation

---

- 1: Initialize cluster centers  $C_k = [l_k, a_k, b_k, x_k, y_k]^T$  by sampling pixels at regular grid steps  $S$ .
- 2: Perturb cluster centers in an  $n \times n$  neighborhood, to the lowest gradient position.
- 3: **repeat** (to the pixel with smallest lab gradient magnitude in 3x3 region)
- 4:   **for** each cluster center  $C_k$  **do**
- 5:     Assign the best matching pixels from a  $2S \times 2S$  square neighborhood around the cluster center according to the distance measure (see next slide)
- 6:   **end for**
- 7:   Compute new cluster centers and residual error  $E$  { $L1$  distance between previous centers and recomputed centers}
- 8: **until**  $E \leq \text{threshold}$

# Notation

<b>N</b>	Number of pixels in the input image
<b>K</b>	Number of Superpixels used to segment the input image
<b>N/K</b>	Approximate size of each superpixel
<b><math>S = \sqrt{N/K}</math></b>	For roughly equally sized superpixels there would be a superpixel centre at every grid interval S

# SLIC Superpixel Segmentation

- Distance function between 2 pixels:

$$D_s = d_{\text{lab}} + \frac{m}{S} d_{x,y}$$

$d_{\text{lab}}$  : lab distance (Euclidean) between the 2 pixels.

$\frac{1}{S} d_{x,y}$  : Euclidean spatial distance, normalized by grid interval  $S$ .

$m$ : **compactness** control of a super pixel. Large values make it more compact.



# Superpixels: More Examples



Approx. 300 to 100 superpixels

# Template Matching

# Template Matching Intro

- Want to find areas of a search image that are similar to given template image  $T$

Template  
Image  $T$



Search Image



Best Matching Patch in Search Image



# General Approaches

- Template-Based:
  - Utilize raw template (pixels) and find best matching patches in search image
    - Sum-of-absolute differences (SAD)
    - Sum-of-squared differences (SSD)
    - Normalized cross-correlation (NCC)

# 1) Sum-of-Absolute Differences (SAD)

- Compute **absolute differences of pixel intensities** of template  $T$  and image patch  $P$  extracted from search image (note that  $P$  is same size as template  $T$ )

$$SAD(P, T) = \sum_{R, G, B} \sum_{x, y} |P(x, y) - T(x, y)|$$

- Compute SAD for all unique patch locations within the search image
- Keep patch with minimum SAD or patches with SAD less than given threshold



# SAD Example

Search Image



Template  
Image  $T$



Negative SAD,  
Origin is in center of patch



# SAD Example

$k^{\text{th}}$  best matching patch



$k = 1$



$k = 10$



$k = 50$



$k = 100$



$k = 150$



$k = 200$




$k = 250$



$k = 500$



$k = 1000$



\*Pixel values scaled between 0 and 1

## 2) Sum-of-Squared Differences (SSD)

- Similar to SAD, but replace absolute differences with **squared differences**

$$SSD(P, T) = \sum_{R, G, B} \sum_{x, y} (P(x, y) - T(x, y))^2$$

- Compute SSD for all unique patches within the search image
- Keep patch with minimum SSD



# SSD Example

Search Image



Template  
Image  $T$



Negative SSD,  
Origin is in center of patch



# SSD Example

$k^{\text{th}}$  best matching patch



$k = 1$



$k = 10$



$k = 50$



$k = 100$



$k = 150$



$k = 200$




$k = 250$



$k = 500$



$k = 1000$



\*Pixel values scaled between 0 and 1

# Illumination Changes

- SAD and SSD can work well if the template and search images have the same brightness
  - Problem: images can have varying illumination conditions

Template  
Image  $T$



Search Image



$k^{\text{th}}$  best matching patch using SSD



$k = 1$     $k = 10$     $k = 50$     $k = 100$



$k = 150$     $k = 200$     $k = 250$     $k = 500$     $k = 1000$

# 3) Normalized Cross-Correlation (NCC)

- Normalize images to remove variations from illumination conditions

$$NCC(P, T) = \sum_{R, G, B} \frac{1}{n-1} \sum_{x, y} \frac{(P(x, y) - \bar{P}) \cdot (T(x, y) - \bar{T})}{\sigma_P \sigma_T}$$

Mean of pixel values in patch  
(each color computed independently)

Standard deviation of pixel values in patch  
(each color computed independently)

*constant ;  
can just be  
calculated once.*

Note: larger values of NCC better!

The maximum value is 1 when two 1-channel signals are exactly the same:

# NCC Example

Search Image



Template  
Image  $T$




NCC,  
Origin is in center of patch






# NCC Example

$k^{\text{th}}$  best matching patch using color images



*Divided NCC by 3  
(for RGB)*



# NCC Example

$k^{\text{th}}$  best matching patch using grayscale images



$k = 1$



$k = 10$



$k = 50$



$k = 100$



$k = 150$



$k = 200$




$k = 250$



$k = 500$




$k = 1000$



# Handling Scale

- Construct fixed-size template of the smallest size you want to detect
- Scan through image pyramid
  - Detects larger scales of object higher in the pyramid
- Efficient scanning method, with less pixels to examine overall
  - Instead of repeatedly scaling template and scanning original full-sized image multiple times





# Summary

- K-Means
  - Choose cluster centers and label every pixel based on its nearest neighbor
  - Minimize total distortion
  - Sensitive to initialization
- Mean-Shift
  - Iteratively shifts data towards peaks
- Superpixel Segmentation
  - Clustering small pixel regions based on color similarity and proximity
- Template Matching
  - SAD, SSD, NCC